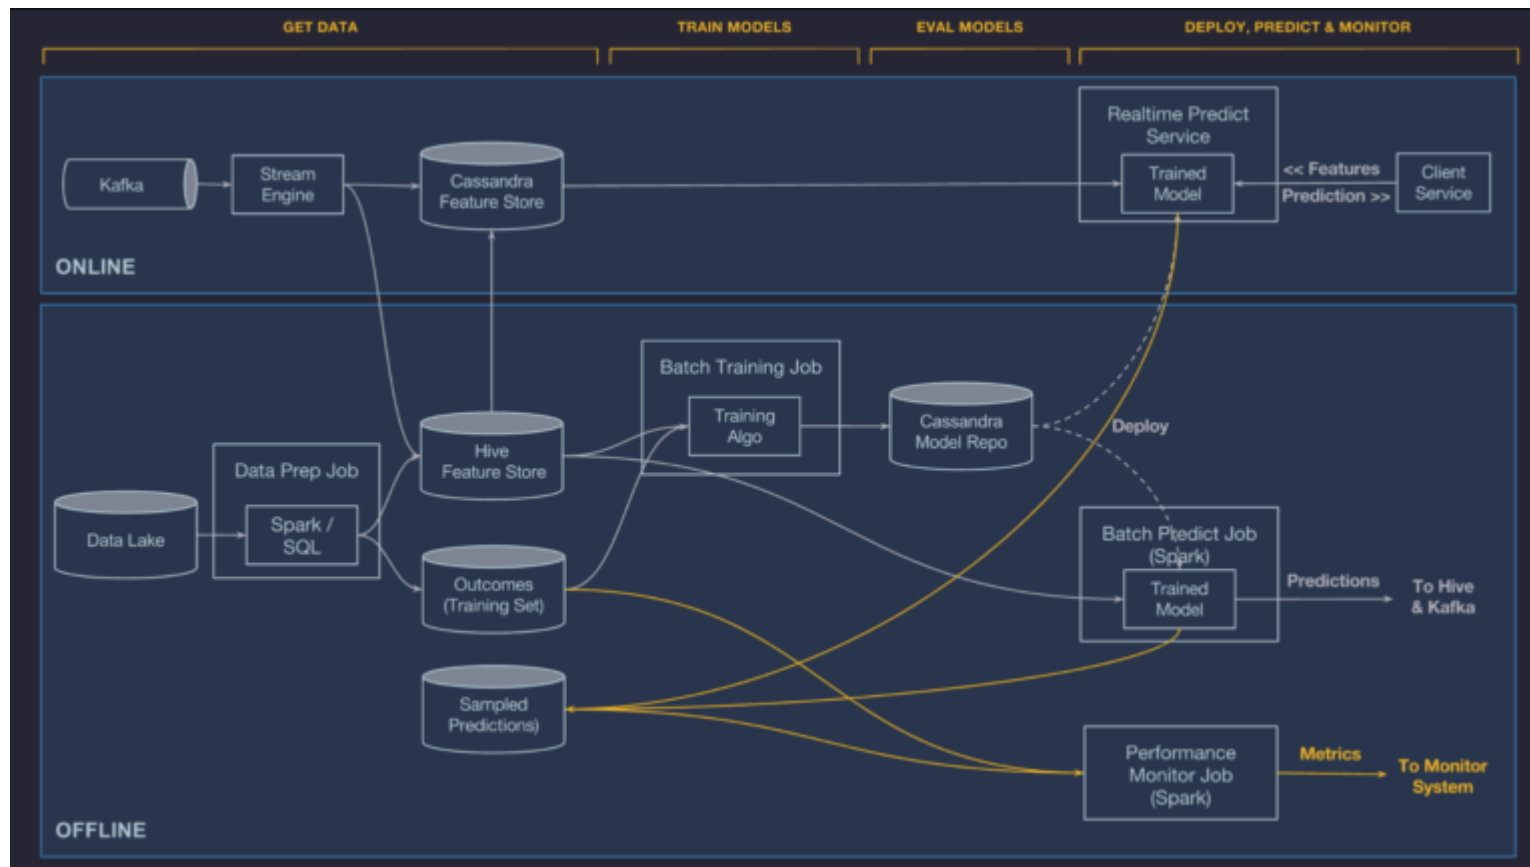


# **Best Practices for Deep Learning**

To support **a viable business** using Deep Learning, you absolutely need an **architecture** that supports sustainable improvement in the presence of **frequent and unexpected changes** in the environment.

“there were no systems in place to build reliable, uniform, and reproducible pipelines for creating and managing training and prediction data at scale.”

--Uber



The architecture supports the following workflow:

- Manage data
- Train models
- Evaluate models
- Deploy, predict and monitor

The Uber system is not a strictly Deep Learning system, but rather a Machine Learning system that can employ many ML methods depending on suitability. It is built on the following open source components: **HDFS**, **Spark**, **Samza**, **Cassandra**, **MLLib**, **XGBoost**, and **TensorFlow**. So, it's a conventional BigData system that incorporates Machine Learning components for its analytics

At the moment, we have approximately 10,000 features in **Feature Store that are used to accelerate machine learning projects**, and teams across the company are adding new ones all the time. Features in the Feature Store are automatically calculated and updated daily.

The model information contains:

- Who trained the model
- Start and end time of the training job
- Full model configuration (features used, hyper-parameter values, etc.)
- Reference to training and test data sets
- Distribution and relative importance of each feature
- Model accuracy metrics
- Standard charts and graphs for each model type (e.g. ROC curve, PR curve, and confusion matrix for a binary classifier)
- Full learned parameters of the model
- Summary statistics for model visualization

“

## TFX: A TensorFlow-based production scale machine learning platform

”

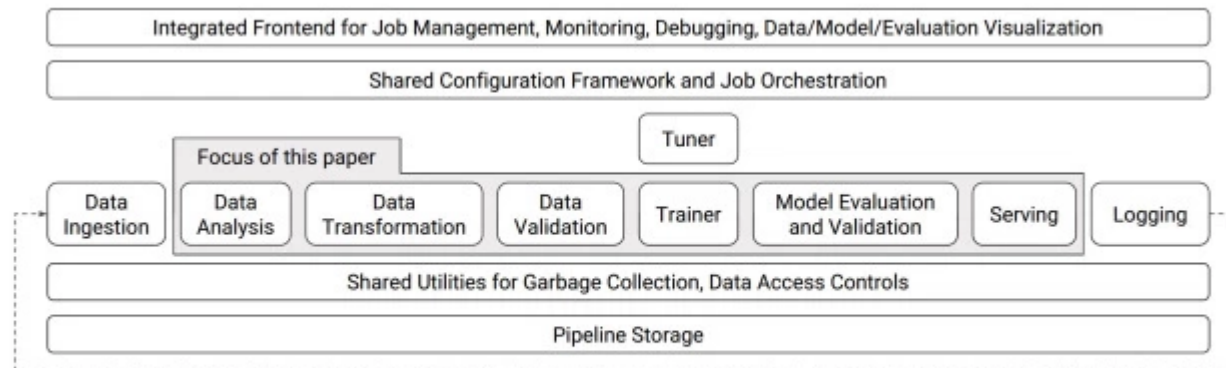


Figure 1: High-level component overview of a machine learning platform.

As compared to available open source DL frameworks, there is a greater emphasis in managing and sharing of meta-information.

# Tensorflow Deploy

## Distributed TensorFlow

how to create a cluster of TensorFlow servers, and how to distribute a computation graph across that cluster.

### Create a cluster

```
tf.train.ClusterSpec({  
    "worker": [  
        "worker0.example.com:2222",  
        "worker1.example.com:2222",  
        "worker2.example.com:2222"  
    ],  
    "ps": [  
        "ps0.example.com:2222",  
        "ps1.example.com:2222"  
    ]  
})
```

/job:worker/task:0

/job:worker/task:1

/job:worker/task:2

/job:ps/task:0

/job:ps/task:1

## Create a server

```
# In task 0:
cluster = tf.train.ClusterSpec({"local": ["localhost:2222", "localhost:2223"]})
server = tf.train.Server(cluster, job_name="local", task_index=0)
```

```
# In task 1:
cluster = tf.train.ClusterSpec({"local": ["localhost:2222", "localhost:2223"]})
server = tf.train.Server(cluster, job_name="local", task_index=1)
```

## Specifying distributed devices

```
with tf.device("/job:ps/task:0"):
    weights_1 = tf.Variable(...)
    biases_1 = tf.Variable(...)

with tf.device("/job:ps/task:1"):
    weights_2 = tf.Variable(...)
    biases_2 = tf.Variable(...)

with tf.device("/job:worker/task:7"):
    input, labels = ...
    layer_1 = tf.nn.relu(tf.matmul(input, weights_1) + biases_1)
    logits = tf.nn.relu(tf.matmul(layer_1, weights_2) + biases_2)
    # ...
    train_op = ...

with tf.Session("grpc://worker7.example.com:2222") as sess:
    for _ in range(10000):
        sess.run(train_op)
```



# TensorFlow Serving

a flexible, high-performance serving system for machine learning models, designed for production environments

TensorFlow Serving 是一个用于机器学习模型 serving 的高性能开源库。它可以将训练好的机器学习模型部署到线上，使用 gRPC 作为接口接受外部调用。更加让人眼前一亮的是，它支持模型热更新与自动模型版本管理。这意味着一旦部署 TensorFlow Serving 后，你再也不需要为线上服务操心，只需要关心你的线下模型训练。


TensorFlow Serving 目前依赖 Google 的开源编译工具 **Bazel**。Bazel 是 Google 内部编译工具 Blaze 的开源版本，功能与性能基本一致。具体的安装可以参考[官方文档](#)。此外还需要安装 **gRPC** (Google 又一个内部工具的开源版)。

```
saver.save(sess, "/tmp/test.ckpt")
```

```
saver.restore(sess, "/tmp/test.ckpt")
```

```
tf.contrib.session_bundle.exporter.Exporter
```

```
model_exporter = exporter.Exporter(saver)
model_exporter.init(
    sess.graph.as_graph_def(),
    named_graph_signatures={
        'inputs': exporter.generic_signature({'x': x}),
        'outputs': exporter.generic_signature({'y': y_pred})
    }
)
model_exporter.export(FLAGS.work_dir,
                      tf.constant(FLAGS.export_version), ses
```



## 模型部署

```
$ bazel build //tensorflow_serving/model_servers:tensorflow_model_server
$ bazel-bin/tensorflow_serving/model_servers/tensorflow_model_server --port=9000 --model_name=test
--model_base_path=/tmp/test/
```

## 客户端

```
from grpc.beta import implementations
import numpy as np
import tensorflow as tf
from tensorflow_serving.apis import predict_pb2
from tensorflow_serving.apis import prediction_service_pb2
tf.app.flags.DEFINE_string('server', 'localhost:9000',
                           'PredictionService host:port')

FLAGS = tf.app.flags.FLAGS
n_samples = 100
host, port = FLAGS.server.split(':')
channel = implementations.insecure_channel(host, int(port))
stub = prediction_service_pb2.beta_create_PredictionService_stub(channel)
```

## Generate test data

```
x_data = np.arange(n_samples, step=1, dtype=np.float32)
x_data = np.reshape(x_data, (n_samples, 1))
```

## Send request

```
request = predict_pb2.PredictRequest()
request.model_spec.name = 'test'
request.inputs['x'].CopyFrom(tf.contrib.util.make_tensor_proto(x_data, shape=[100, 1]))
result = stub.Predict(request, 10.0) # 10 secs timeout
```

配置一下 bazel 的 BUILD 文件

```
bazel build //tensorflow_serving/test:test_client && ./bazel-bin/tensorflow_serving/test/test_client
```